

Human Language Technologies

Twitter Sentiment Analysis

Gianmarco Ricciarelli
gianmarcoricciarelli@gmail.com
Project's Repository: GitHub

Abstract

Nowadays, when delving into the Social Network landscape, a user can choose among different paths, depending on the type of experience he/she is searching. Regardless the type of platform that is chosen by the user, either Facebook, or Instagram or Twitter, the amount of textual data that is produced every day is massive. With this report, I describe the project I developed for the *Human Language Technologies* class hosted by the University of Pisa's Master's Degree in Computer Science, that is, a prediction-oriented analysis of a dataset composed by almost 40000 labeled tweets via a series of models like SVM, NB, CNN and LSTM.

1 Corpus Collection

The amount of textual data produced by the users on the various Social Networks on a day-by-day basis is massive, and, for this reason, it is not too difficult to scrape the web in order to build a corpus composed by an acceptable amount of documents. Among all the available platforms, I choosed to work with Twitter, since the documents, that is, the tweets, produced by the users are limited to be 280 characters long at most, and since the API that is used for communicating with the server is very simple and intuitive to use. In order to properly train the models that I wanted to compare, I assembled a corpus of almost 40000 labeled tweets by requesting Twitter to download the documents used for the *International Workshop on Semantic Evaluation* competitions from 2013 to 2017. I choosed to work with the documents provided by the SemEval competitions because every tweet was hand-labeled by a human, that is, guaranteeing a sound catalogation of the sentiments expressed by the tweets. A tweet can be labeled either as positive, negative, or neutral. Despite the tweets' labels being so well defined, a SemEval dataset is quite small if it is taken

on its own. This is due to the fact that, since Twitter is a very dynamic and constantly upgraded platform, some of the documents provided by each dataset were deleted by the users. This is why I decided to merge the datasets provided by the various competitions in order to obtain the final version of the corpus.

2 Corpus Analysis

The original version of the corpus I assembled was composed by more than 40000 labeled tweets. As I said before, some of the tweets were deleted by the users, resulting in a document containing the 'Not Available' string. Moreover, the corpus contained a small percentage of duplicate tweets. By removing the not available tweets and the duplicates, I obtained the final version of the corpus composed by 39308 labeled documents. In particular, the obtained corpus contains:

- 15092 positive tweets;
- 6007 negative tweets;
- 18209 neutral tweets;

As we can see, the majority of the documents composing the corpus are labeled as neutral, and also the positive tweets are very well represented. Conversely, the negative ones are poorly represented, and, as we will see, this fact will be crucial for their classification in the learning and testing phases of the algorithms I will apply. The type of subdivision used for the documents allows the corpus to be considered as containing both subjective and objective tweets. By considering this new subdivision, we can see that the corpus contains:

- 21099 subjective tweets;

- 18209 objective tweets;

This time, both the two sets of documents are well represented. In order to provide a richer analysis, I decided to apply the classification algorithms on three types of problems; the positive-vs-negative-vs-neutral problem, the positive-vs-negative problem and finally the subjective-vs-objective problem. For starting my analysis, I tokenized the tweets contained in the corpus and I observed the distribution of the word frequencies.




Figure 1: Word Frequencies Distribution.

As we can see in Figure 1, the words' frequencies follow the well known Zipf's law, with a small number of words that appears frequently, the so called stop words, and the majority of the remaining words that appears less and less frequently. Following the analysis proposed in [1], I updated the tokens obtained by the early stage of the analysis by adding the tags from the application of a Part of Speech Tagging procedure, and I observed the tags' distribution between the corpus' documents when considering only the positive/negative tweets and the subjective/objective tweets.

As I said before, [1] proposes the following formula to perform a pairwise comparison of tags distributions for each tag and two sets:

$$P_{1,2}^T = \frac{N_1^T - N_2^T}{N_1^T + N_2^T}$$

where N_1^T and N_2^T are numbers of tag T occurrences in the first and second sets respectively. As we can see from Figure 2, POS tags for positive and negative




Figure 2: Tag Distribution for Positive and Negative tweets.

tweets are not distributed evenly, hence we can infer the sentiment behind a document by looking at its tags. We can see that the presence of a plural proper noun (NNPS) is a strong indicator for the positive label as well as past participle verbs (VBN) and modal (MD). For the negative label, we can see that 'to' (TO) is a strong indicator, as well as symbol (SYM) and singular proper nouns (NNP).




Figure 3: Tag Distribution for Objective and Subjective tweets.

In Figure 3 we can observe the tags' distribution for the subjective and objective tweets. We can see that possessive wh-pronouns (WPS), possessive pronouns (PRP\$) and superlative adjectives (JJS) are all indicators for objective tweets, while 'to' (TO), foreign words (FW) and wh-determiners (WDT) are indicators for subjective tweets. Finally, I have collected some of the most used words for positive and negative tweets and represented them via wordclouds, in order

to provide a better understanding of the sentiments that labels the documents composing the corpus, and for providing some visual representation of the informations showed with Figure 2 and Figure 3.

Figure 4: Positive tokens wordcloud.



Figure 5: Negative tokens wordcloud

The Human Language Technologies research’s field is very popular at the moment, and, in particular, the Sentiment Analysis branch has a rich literature and many tools and resources available for learning. Having analysed the obtained corpus, as described in Section 1 and Section 2, the next step for my research was to implement a Classification procedure, in which I decided to apply some popular Machine Learning models like Support Vector Machines, Naïve Bayes, Convolutional Neural Networks and Long Short-term Memory Recurrent Neural Networks, as described in [2, 3, 4, 5, 6]. Being the corpus composed of documents labeled with three distinct labels, that is, posi-

- Classification of positive tweets vs negative tweets;
- Classification of positive tweets vs negative tweets vs neutral tweets;
- Classification of subjective tweets vs objective tweets;

3.1 Preprocessing Routine

1. Tokenization: the original documents are split into tokens via the TweetTokenizer tokenizer provided by NLTK. This tokenizer is specifically trained to work on tweets;
2. Part Of Speech Tagging: each token list is tagged via the application of a POS tagger, also provided by NLTK;
3. Cleaning: for each token list, the token that are not alphanumeric and that contains stop words are removed;
4. Lemmatization: each one of the token that came out of the Cleaning phase is transformed in its original lemma;

As an example, if we apply the Preprocessing Routine to the document:

```
April has to live. Sharknado wouldn't
be the same without her tiny chainsaw.
#aprillives,
```

we obtain the token list:

```
['april', 'live', 'sharknado', 'without',
'tiny', 'chainsaw'].
```

3.2 Pipeline

Before feeding the corpus to the Naïve Bayes and Support Vector Machine classifiers I further preprocessed the documents via the application of a Pipeline composed by:

1. A Vectorizer, that is, the CountVectorizer provided by scikit-learn, which takes in input the tokenized documents from the Preprocessing Routine and, for each tweet, returns a list of integers where each integer corresponds to a word in the tweet;
2. A Transformer, that is, the TfidfTransformer provided by scikit-learn, which takes the matrix created by the Vectorizer and transform it to a normalized or tf-idf representation;
3. A Selectioner, that is, the SelectKBest provided by scikit-learn, which takes the output of the Transformer and select the k most representative features via the application of the chi-squared test;

Moreover, for the Selectioner, the number of features selected, that is, the k parameter, ranges between 1000 and 5000. This parameter is adapted depending on the task during the Validation Routine, which is described in the next section.

3.3 Validation Routine

Every Machine Learning algorithm is characterized by a number of parameters, that is, the so-called hyperparameters, that the user can modify accordingly to the type of task he/she is tackling. Being the task of tuning every single hyperparameter a cumbersome one, I decided to apply the well-know Random Grid

Search technique, as described in [7], in order to discover the most effective combination of hyperparameters for the models I want to compare. I used the 80% of the corpus as a training set, while the remaining 20% was used as testing set. Moreover, since the training/validation set extracted from the corpus presents a reasonable amount of examples, for each "point" produced by the Random Grid Search, I applied the K-Fold Cross Validation technique, setting the value for the K parameter to 3. Moreover, since the performance of algorithms like the Convolutional Neural Networks and the Long Short-term Memory Recurrent Neural Networks are greatly influenced by the network's architecture, I decided, in order to ease the validation process for each one of the two algorithms, to fix a simple architecture and to apply the validation techniques to that fixed structure for obtaining the best results. The architecture I selected for the CNN is composed by an Embedding Layer, followed by a Convolutional Layer, a Max Pooling and finally by a Fully Connected Neural Network. The architecture I selected for the LSTM is composed by an Embedding Layer, followed by a LSTM Layer and finally by a Fully Connected Neural Network. For enriching my analysis of the models' performances, I used and validated three types of word embeddings for the CNN and the LSTM algorithms:

- a word embedding learned during the training phase with 100 entries for each word;
- the GloVe [8] pre-trained word vectors with 100 entries for each word;
- the fastText [9] pre-trained word vectors with 300 entries for each word;

For now on, I will refer to the CNN/LSTM using the word embeddings learned during the training phase with the subscript A, while I will use the subscript G for the GloVe pre-trained word vectors and F for the fastText pre-trained word vectors, respectively.

3.4 Positive vs Negative Classification

As we might think, this is the easier one among the tasks I have decided to approach. This is because the original multilabel classification problem is downgraded to a binary classification problem, and, moreover, by removing the documents labeled as neutral we obtain a corpus where the positive labeled tweets

are very well separated from the negative ones, hence resulting in a easier classification from the Machine Learning point of view.

Classifier	Score
<i>SVM</i>	0.85
<i>NB</i>	0.82
<i>CNN_A</i>	0.83
<i>CNN_G</i>	0.81
<i>CNN_F</i>	0.85
<i>LSTM_A</i>	0.83
<i>LSTM_G</i>	0.84
<i>LSTM_F</i>	0.83

Table 1: Models’ comparison for the positive-vs-negative classification task.

In Table 1 we can see the results in terms of Accuracy for the application of the models to the corpus containing only positive and negative documents. As we can see, in general, the Accuracy is above 0.80, this due to the fact that, as I said, this classification task is not so difficult compared to the other two. The higher Accuracy is reached by the SVM model and by the CNN using the fastText pre-trained word vectors, while the worst Accuracy is obtained by the CNN using the GloVe pre-trained word vectors.

3.5 Positive vs Negative vs Neutral Classification

In order to takle this task, the entire corpus is used, with the only modification to the documents being the application of the preprocessing routine described in Section 3.1. This type of task results to be more difficult compared to the one described in Section 3.4, both in term of time and computational resources.

Classifier	Score
<i>SVM</i>	0.64
<i>NB</i>	0.61
<i>CNN_A</i>	0.64
<i>CNN_G</i>	0.58
<i>CNN_F</i>	0.62
<i>LSTM_A</i>	0.61
<i>LSTM_G</i>	0.62
<i>LSTM_F</i>	0.64

Table 2: Models’ comparison for the positive-vs-negative-vs-neutral classification task.

In Table 2 we can see the results in terms of Accuracy for the application of the models to the entire corpus. As we can see, in general, the Accuracy tend to be greater than 0.60, with the exception of the CNN model using the GloVe pre-trained word vectors. This time the best models are, respectively, the SVM model, the CNN using word embeddings learned during the training phase and finally the LSTM model using the fastText pre-trained word vectors.

3.6 Subjective vs Objective Classification

Finally I discuss the last of the tasks that I takled for my project, that is, the classification of subjective and objective tweets. As for the problem described in Section 3.5, also for this type of classification the entire corpus was used. For the tasks described in Sections 3.4 and 3.5 the labels were substituted with positive integers from 0 to 1 for the positive-vs-negative classification, and from 0 to 2 for the positive-vs-negative-vs-neutral classification. For this type of classification I labeled all the positive and negative documents, that is, the subjective ones, with the integer 0, while the neutral ones, that is, the objective ones, were labeled with the integer 1. As for the task described in Section 3.5, also for this task the request in terms of time and computational resources was high.

Classifier	Score
<i>SVM</i>	0.67
<i>NB</i>	0.64
<i>CNN_A</i>	0.64
<i>CNN_G</i>	0.64
<i>CNN_F</i>	0.66
<i>LSTM_A</i>	0.68
<i>LSTM_G</i>	0.66
<i>LSTM_F</i>	0.68

Table 3: Models’ comparison for the subjective-vs-objective classification task.

In Table 3 we can see the results in terms of Accuracy for the application of the models to the entire corpus composed by subjective and objective tweets. By observing the Table we can notice that, despite the fact that also for this task the entire corpus was used, the average Accuracy is greater than the one obtained for the task described in Section 3.5. This is due to the fact that this type of task is a binary classification task, and not a multilabel classification one, in which each one of the two labels is well represented, as we

can see in Section 2. We can see that the best models are both LSTMs, one using the fastText pre-trained word vectors and the other using a word embedding learned during the training phase, followed closely by the SVM model.

4 Conclusions

The Sentiment Analysis research field is very active at the moment, with a great number of paper available and many tools and resources that can be used for learning and enriching our knowledge about the subject. The possibility to freely access great and constantly updated amount of data, like, for example, Twitter’s tweets, Facebook’s posts, screenplays, etc. etc., makes it easy to experiment with the available models or even to create new ones in order to obtain better performances in comparison with the old ones. With this project I learned about how to obtain a good amount of documents from a reliable source, which in this case is represented by the *International Workshop on Semantic Evaluation*, how to analyze textual data in order to discover hidden relationships between the documents composing a corpus, and, most importantly, I learned how to use well-know models like Support Vector Machines, Naïve Bayes, Convolutional Neural Networks and Long Short-term Memory Recurrent Neural Networks in order to classify the sentiment behind a document written by an human user, which nowadays can be an useful resource to be used in a variety of application’s fields. I think that the performances I obtained in term of Accuracy score, especially the CNN and LSTM ones, might be upgraded via the use of different architectures, but, as I said before, for my project I preferred a simpler approach.

References

- [1] Pak, Alexander & Paroubek, Patrick. (2010). Twitter as a Corpus for Sentiment Analysis and Opinion Mining. *Proceedings of LREC*. 10.
- [2] Nakov, P. & Rosenthal, S. & Kiritchenko, S. et al. Developing a successful SemEval task in sentiment analysis of Twitter and other social media texts. *Lang Resources & Evaluation* 50, 35–65 (2016). <https://doi.org/10.1007/s10579-015-9328-1>
- [3] Mohammad, Saif & Kiritchenko, Svetlana & Zhu, Xiaodan. NRC-Canada: Building the State-of-the-Art in Sentiment Analysis of Tweets. *Second Joint Conference on Lexical and Computational Semantics, Volume 2: Proceedings of the Seventh International Workshop on Semantic Evaluation*, June 2013, Atlanta, Georgia, USA.
- [4] Nakov, Preslav & Ritter, Alan & Rosenthal, Sara & Sebastiani, Fabrizio & Stoyanov, Veselin. SemEval-2016 Task 4: Sentiment Analysis in Twitter. *Proceedings of the 10th International Workshop on Semantic Evaluation (SemEval-2016)*, June 2016, San Diego, California.
- [5] Esuli, Andrea. ISTI-CNR at SemEval-2016 Task 4: Quantification on an Ordinal Scale. *Proceedings of the 10th International Workshop on Semantic Evaluation (SemEval-2016)*, June 2016. Association for Computational Linguistics, San Diego, California.
- [6] Attardi, Giuseppe & Sartiano, Daniele. (2016). UniPI at SemEval-2016 Task 4: Convolutional Neural Networks for Sentiment Classification. 220–224. 10.18653/v1/S16-1033.
- [7] James Bergstra and Yoshua Bengio. 2012. Random search for hyper-parameter optimization. *J. Mach. Learn. Res.* 13, 1 (February 2012), 281–305.
- [8] Jeffrey Pennington & Richard Socher & Christopher D. Manning. GloVe: Global Vectors for Word Representation. *Empirical Methods in Natural Language Processing (EMNLP)*, 2014, 1532–1543 <http://www.aclweb.org/anthology/D14-1162>
- [9] Mikolov, Tomas & Grave, Edouard & Bojanowski, Piotr & Puhersch, Christian & Joulin, Armand. Advances in Pre-Training Distributed Word Representations. *Proceedings of the International Conference on Language Resources and Evaluation (LREC 2018)*, 2018.